

Linear & NonLinear Project

Adham Suliman

July 17, 2018

1. & 2) & 3)

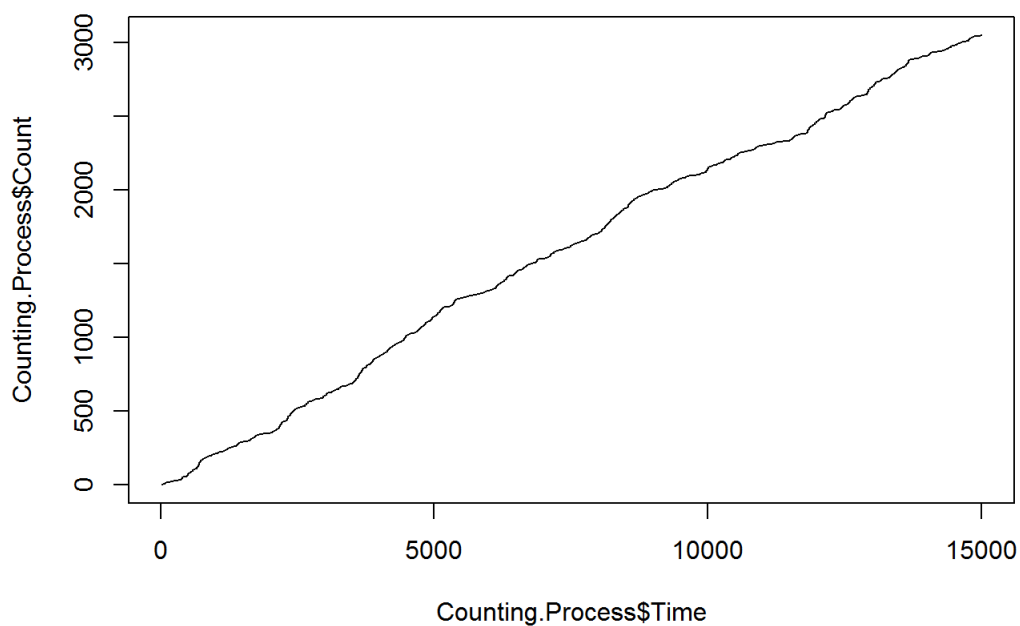
```
Course.Project.Data <- read.csv('LinearNonLinear_MalfunctionData.csv', header = T)
Counting.Process<-as.data.frame(cbind(Time=Course.Project.Data$Time,Count=1:length(Course.Project.Data$Time)
))
Counting.Process[1:20,]
```

##	Time	Count
## 1	18.08567	1
## 2	28.74417	2
## 3	34.23941	3
## 4	36.87944	4
## 5	37.84399	5
## 6	41.37885	6
## 7	45.19283	7
## 8	60.94242	8
## 9	66.33539	9
## 10	69.95667	10
## 11	76.17420	11
## 12	80.48524	12
## 13	81.29133	13
## 14	86.18149	14
## 15	91.28642	15
## 16	91.75162	16
## 17	98.29452	17
## 18	142.58741	18
## 19	149.82484	19
## 20	151.58587	20

```
max(Counting.Process$Time)
```

```
## [1] 14999.94
```

```
plot(Counting.Process$Time,Counting.Process$Count,type="s")
```




The counting process

trajectory looks pretty smooth and grows steadily.

What does it tell you about the character of malfunctions and the reasons causing them? This tells me that as time goes on, there seems to be a pretty constant count which could mean an average rate or “intensity” of the issue occurring over time.

3.1)

```
plot(Counting.Process$Time, Counting.Process$Count/Counting.Process$Time, type="l", ylab="Cumulative Intensity")
abline(h=Counting.Process$Count[length(Counting.Process$Count)]/
       Counting.Process$Time[length(Counting.Process$Time)])
abline(h=mean(Counting.Process$Count/Counting.Process$Time))
```

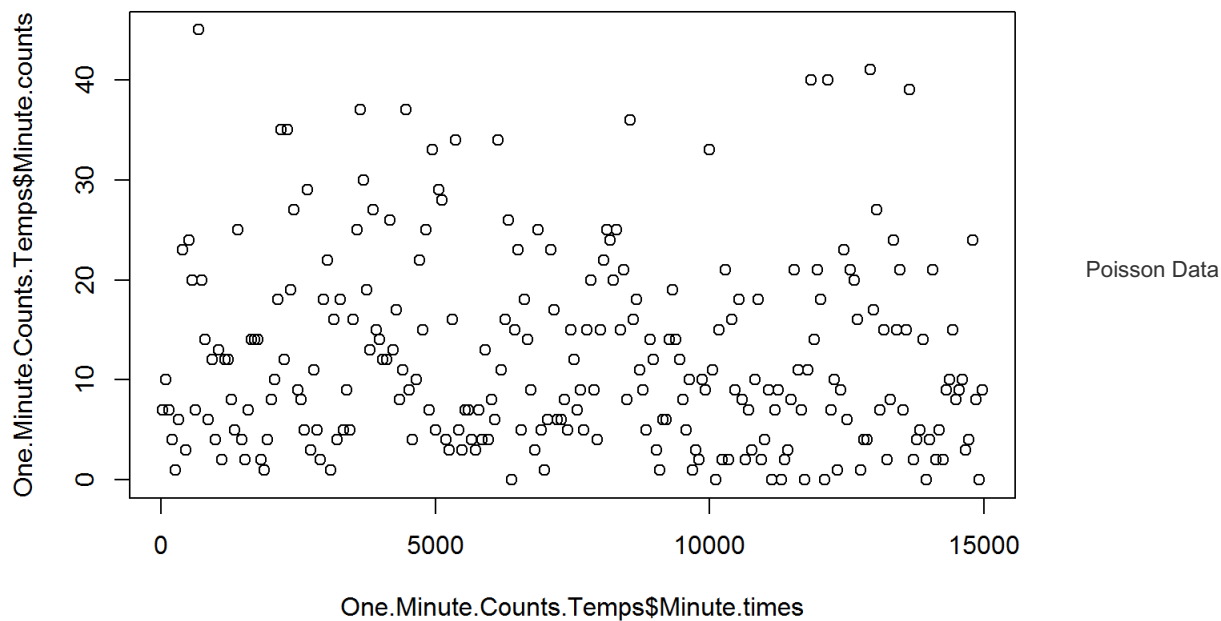


```
c(Last.Intensity=Counting.Process$Count[length(Counting.Process$Count)]/
   Counting.Process$Time[length(Counting.Process$Time)],
   Mean.Intensity=mean(Counting.Process$Count/Counting.Process$Time))
```

```
## Last.Intensity Mean.Intensity
##      0.2036008      0.2095305
```

4.

```
Course.Project.Data$Minute <- floor(Course.Project.Data$Time/60)
Course.Project.Data1 <-
  Course.Project.Data%>%
    dplyr::select(Minute, Temperature)%>%
    dplyr::group_by(Minute) %>%
    dplyr::summarise(Minute.counts = c(n()), Miunute.Temps = c(mean(Temperature)))
One.Minute.Counts.Temps<- Course.Project.Data1%>%
    dplyr::mutate(Minute.times=c((Minute*60)+30))%>%
    dplyr::select(Minute.times, Minute.counts, Miunute.Temps)
NCourse.Project.Data <-vector()
for (x in 1:250){
  NCourse.Project.Data[x]<-(30+(60*(x-1)))
}
NCourse.Project.Data <- as.data.frame(matrix(NCourse.Project.Data,ncol=1),col.names="Minute.times")
colnames(NCourse.Project.Data)="Minute.times"
One.Minute.Counts.Temps <-
  NCourse.Project.Data %>%
    left_join(One.Minute.Counts.Temps,by="Minute.times")
One.Minute.Counts.Temps$Minute.counts[is.na(One.Minute.Counts.Temps$Minute.counts)] <- 0
plot(One.Minute.Counts.Temps$Minute.times,One.Minute.Counts.Temps$Minute.counts)
```



```
Test.Deviance.Overdispersion.Poisson<-function(Sample.Size,Parameter.Lambda){
  my.Sample<-rpois(Sample.Size,Parameter.Lambda)
  Model<-glm(my.Sample~1,family=poisson)
  Dev<-Model$deviance
  Deg.Fred<-Model$df.residual
  (((Dev/Deg.Fred-1)/sqrt(2/Deg.Fred))>-1.96)&(((Dev/Deg.Fred-1)/sqrt(2/Deg.Fred))<=1.96))*1
}
Test.Deviance.Overdispersion.Poisson(100,1)
```

```
## [1] 1
```

```
sum(replicate(300,Test.Deviance.Overdispersion.Poisson(100,1)))
```

```
## [1] 272
```

```
exp(glm(rpois(1000,2)~1,family=poisson)$coeff)
```

```
## (Intercept)
##          1.947
```

Negative Binomial Data

```
Test.Deviance.Overdispersion.NBinom<-function(Sample.Size,Parameter.prob){
  my.Sample<-rnbinom(Sample.Size,2,Parameter.prob)
  Model<-glm(my.Sample~1,family=poisson)
  Dev<-Model$deviance
  Deg.Fred<-Model$df.residual
  (((Dev/Deg.Fred-1)/sqrt(2/Deg.Fred))>-1.96)&(((Dev/Deg.Fred-1)/sqrt(2/Deg.Fred))<=1.96))*1
}
sum(replicate(300,Test.Deviance.Overdispersion.NBinom(100,.2)))
```

```
## [1] 0
```

Apply the test to the one minute event counts Do you see signs of over-dispersion? Alpha is greater than 0 due to the significant p-value below which shows over dispersion.

```
GLM.model<-glm(One.Minute.Counts.Temps$Minute.counts~1,family=poisson)
GLM.model
```

```
##
## Call: glm(formula = One.Minute.Counts.Temps$Minute.counts ~ 1, family = poisson)
##
## Coefficients:
## (Intercept)
## 2.503
##
## Degrees of Freedom: 249 Total (i.e. Null); 249 Residual
## Null Deviance: 1799
## Residual Deviance: 1799 AIC: 2789
```

```
(Disp.test <- dispersiontest(GLM.model))
```

```
##
## Overdispersion test
##
## data: GLM.model
## z = 8.5747, p-value < 2.2e-16
## alternative hypothesis: true dispersion is greater than 1
## sample estimates:
## dispersion
## 7.377975
```

Q: Does the test show overdispersion? A: Yes it does because the alpha is above 0 and pour p-value is significant which shows over dispersion.

4.1.3 Test against negative binomial distribution Definitely not an NB model.

```
GLM.NM<-glm.nb(Minute.counts~1,One.Minute.Counts.Temps)
summary(GLM.NM)
```

```
##
## Call:
## glm.nb(formula = Minute.counts ~ 1, data = One.Minute.Counts.Temps,
## init.theta = 1.747516571, link = log)
##
## Deviance Residuals:
## Min 1Q Median 3Q Max
## -2.6951 -0.9389 -0.2977 0.4958 2.0931
##
## Coefficients:
## Estimate Std. Error z value Pr(>|z|)
## (Intercept) 2.50275 0.05115 48.93 <2e-16 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for Negative Binomial(1.7475) family taken to be 1)
##
## Null deviance: 278.5 on 249 degrees of freedom
## Residual deviance: 278.5 on 249 degrees of freedom
## AIC: 1746.7
##
## Number of Fisher Scoring iterations: 1
##
## Theta: 1.748
## Std. Err.: 0.179
##
## 2 x log-likelihood: -1742.697
```


```
odTest(GLM.NM)
```

```
## Likelihood ratio test of H0: Poisson, as restricted NB model:
## n.b., the distribution of the test-statistic under H0 is non-standard
## e.g., see help(odTest) for details/references
##
## Critical value of test statistic at the alpha= 0.05 level: 2.7055
## Chi-Square Test Statistic = 1044.585 p-value = < 2.2e-16
```

Does this test show overdispersion? Yes it does due to the significant P-value.

5. Find the distribution of Poisson Intensity 5.1 Kolmogorov-Smirnov Test

```
sample1=rnorm(100)
sample2=rnorm(100,1,2)
Cum.Distr.Functions <- data.frame(sample1,sample2)
ecdfplot(~ sample1 + sample2, data=Cum.Distr.Functions, auto.key=list(space='right'))
```



```
ks.test(sample1,sample2)
```

```
##
## Two-sample Kolmogorov-Smirnov test
##
## data: sample1 and sample2
## D = 0.48, p-value = 1.972e-10
## alternative hypothesis: two-sided
```

What does this output tell you about equivalence of the two distributions? Because the p-value is significant, these samples have come from different populations.

```
ks.test(sample1,"pnorm",mean=0,sd=1)
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: sample1
## D = 0.072997, p-value = 0.6609
## alternative hypothesis: two-sided
```

What does this output tell you? Because the p-value is not significant, the 1st samples does come from a population with a normal distribution $N(0,1)$.

```
ks.test(sample2,"pnorm",mean=0,sd=1)
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: sample2
## D = 0.45719, p-value < 2.2e-16
## alternative hypothesis: two-sided
```

Check equivalence of the empirical distribution of sample2 and theoretical distribution Norm(0,1). The second sample comes from a population that doesn't have a normal distribution N(0,1).

Apply Kolmogorov-Smirnov test to Counting.Process\$Time and theoretical exponential distribution with parameter equal to average intensity.

```
Time.diff <- diff(Counting.Process$Time)
intensity <- 1:length(Counting.Process$Time) / (Counting.Process$Time)
mean.intensity <- mean(intensity)
KS.Test.Event.Interval <- ks.test(Time.diff, "pexp", rate=mean.intensity)
c(KS.Test.Event.Interval$statistic, p.value=KS.Test.Event.Interval$p.value)
```

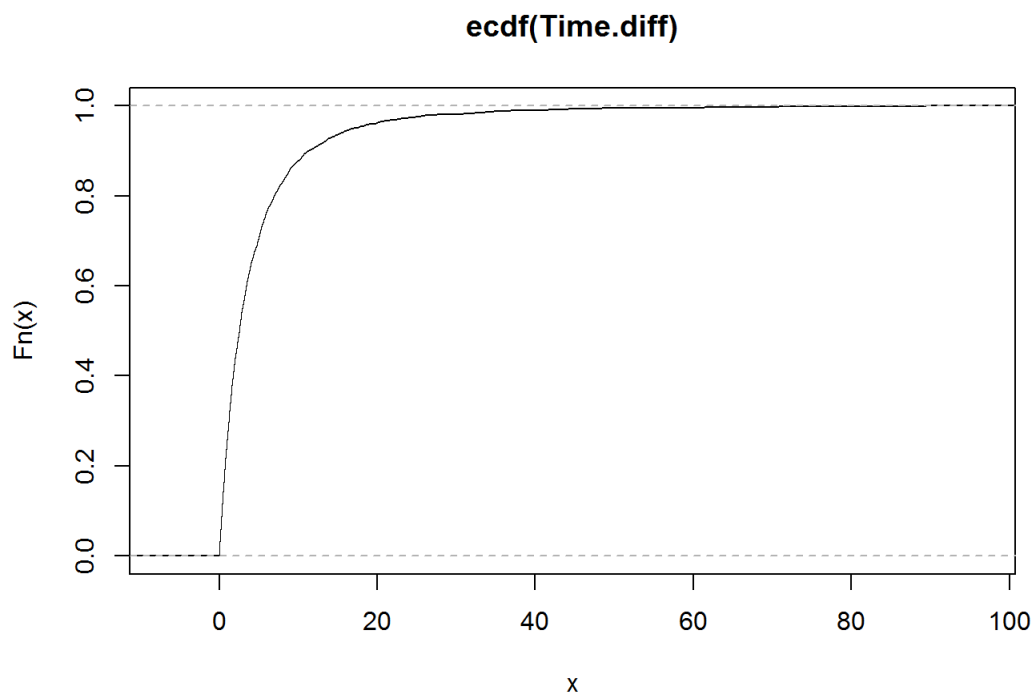
```
##          D    p.value
## 0.0950615 0.0000000
```

theoretical exponential distribution with parameter equal to average intensity.

```
Time.diff <- diff(Counting.Process$Time)
ks.test(Time.diff, "pexp", mean(Time.diff), sd(Time.diff))
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: Time.diff
## D = 0.78677, p-value < 2.2e-16
## alternative hypothesis: two-sided
```

```
plot(ecdf(Time.diff))
```




```
order(One.Minute.Counts.Temps$Minute.counts, decreasing=T)
```

```
## [1] 12 216 198 203 228 61 75 143 37 39 90 103 83 167 62 45 85 86 41 65 218 70 106
## [24] 24 60 81 115 136 139 9 137 223 247 7 109 119 208 51 79 135 141 172 193 200 210 225
## [47] 235 10 13 131 138 211 40 63 156 36 50 55 111 145 176 182 201 72 120 217 53 59 89
## [70] 105 144 174 212 66 80 108 125 130 134 140 170 220 224 227 241 14 28 29 30 67 112 149
## [93] 155 157 199 232 18 64 71 99 16 20 21 38 68 69 126 150 158 47 74 104 146 168 194
## [116] 197 2 35 78 161 165 181 205 240 244 42 57 76 113 128 132 147 166 175 185 188 207 239
## [139] 243 250 22 34 43 73 101 123 142 159 177 192 222 242 248 1 3 11 27 82 93 94 97
## [162] 127 179 187 195 204 219 226 6 15 102 118 121 122 153 154 209 23 44 48 56 58 84 91
## [185] 110 116 124 129 148 160 231 237 4 17 25 33 54 77 87 95 98 100 133 184 214 215 230
## [208] 234 246 8 46 88 92 96 114 151 163 180 191 245 19 26 31 49 164 171 173 178 183 190
## [231] 221 229 236 238 5 32 52 117 152 162 206 213 107 169 186 189 196 202 233 249
```

```
Event.Intensities <- One.Minute.Counts.Temps$Minute.counts/60
hist(Event.Intensities)
```

Histogram of Event.Intensities



What distribution does this

histogram remind you of? This distributino reminds me of a poissoin distribution.

```
start <- seq(0,5,.2)
(Fitting.Normal <- fitdistr(Event.Intensities,densfun="normal"))
```

```
##      mean      sd
## 0.203600000 0.158227459
## (0.010007183) (0.007076147)
```

```
(Fitting.Exponential <- fitdistr(Event.Intensities, densfun="exponential"))
```

```
##      rate
## 4.9115914
## (0.3106363)
```

Logs

```
Event.Intensitiesl <- Event.Intensities + 10e-15
(Fitting.Logs<-fitdistr(Event.Intensitiesl,densfun="lognormal"))
```

```
##      meanlog      sdlog
## -2.8474770 5.4111039
## ( 0.3422283) ( 0.2419919)
```

```
Fitting.Logs
```

```
##      meanlog      sdlog
##    -2.8474770    5.4111039
##    ( 0.3422283) ( 0.2419919)
```

```
KS.LogNormal <- ks.test(Event.Intensities1,"plnorm",meanlog= -2.8474770 ,sd= 5.4111039)
```

```
## Warning in ks.test(Event.Intensities1, "plnorm", meanlog = -2.847477, sd = 5.4111039): ties should
## not be present for the Kolmogorov-Smirnov test
```

```
KS.LogNormal
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: Event.Intensities1
## D = 0.39525, p-value < 2.2e-16
## alternative hypothesis: two-sided
```

Beta

```
mu <- mean(Event.Intensities)
var <- var(Event.Intensities)
estBetaParams <- function(mu, var) {
  alpha <- ((1 - mu) / var - 1 / mu) * mu ^ 2
  beta <- alpha * (1 / mu - 1)
  return(params = list(alpha = alpha, beta = beta))
}
x <- estBetaParams(mu,var)
x[2]
```

```
## $beta
## [1] 4.340912
```

```
KS.beta <- ks.test(Event.Intensities,"pbeta", shape1=1.109756,shape2=4.340912)
```

```
## Warning in ks.test(Event.Intensities, "pbeta", shape1 = 1.109756, shape2 = 4.340912): ties should
## not be present for the Kolmogorov-Smirnov test
```

Exponential distribution should be closer but it's not.

```
KS.Normal <- ks.test(Event.Intensities,"pnorm",mean=Fitting.Normal$estimate[1] ,sd= Fitting.Normal$estimate[
2])
```

```
## Warning in ks.test(Event.Intensities, "pnorm", mean = Fitting.Normal$estimate[1], : ties should not
## be present for the Kolmogorov-Smirnov test
```

```
names(Fitting.Exponential)
```

```
## [1] "estimate" "sd"      "vcov"     "n"        "loglik"
```

```
KS.Exp <- ks.test(Event.Intensities,"pexp",rate=Fitting.Exponential$estimate[1])
```

```
## Warning in ks.test(Event.Intensities, "pexp", rate = Fitting.Exponential$estimate[1]): ties should
## not be present for the Kolmogorov-Smirnov test
```

```
c(KS.Normal$statistic,P.Value=KS.Normal$p.value)
```

```
##           D           P.Value
## 0.1326020316 0.0003039941
```



```
c(KS.Exp$statistic,P.Value=KS.Exp$p.value)
```

```
##           D           P.Value
## 0.115233049 0.002615812
```

What do you conclude from these tests? These tests are telling that the data doesn't fit a normal distribution nor an exponential distribution because of the significant p values.

Try to fit gamma distribution directly using `fitdistr()`

```
#Fitting.Gamma <- fitdistr(Event.Intensities,densfun="gamma")
```

NANs are produced, but why? Since there is a negative value within the data, the gamma distribution cannot be utilized because it can take only positive values.

Intensity is gamma with a normal distribution.

```
first.moment <- mean(Event.Intensities)
second.moment <- var(Event.Intensities)*(length(Event.Intensities))/(length(Event.Intensities)+1)
beta <- first.moment/second.moment
alpha <- first.moment*beta
(Moments.Rate <- beta)
```

```
## [1] 8.132182
```

```
(Moments.Shape <- alpha)
```

```
## [1] 1.655712
```

```
KS.Test.Moments.Gamma <- ks.test(Event.Intensities,"pgamma",rate=beta, shape=alpha)
```

```
## Warning in ks.test(Event.Intensities, "pgamma", rate = beta, shape = alpha): ties should not be
## present for the Kolmogorov-Smirnov test
```

```
KS.Test.Moments.Gamma
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: Event.Intensities
## D = 0.057807, p-value = 0.3737
## alternative hypothesis: two-sided
```

What distribution for the one-minute intensity of malfunctions do you choose? In the end, I would choose the beta distribution due to its high P value.

What distribution of one-minute malfunctions counts follow from your choice? If I couldn't use the beta distribution, I would choose the gamma distribution.

```
rbind(KS.Moments=c(KS.Test.Moments.Gamma$statistic,P.Value=KS.Test.Moments.Gamma$p.value),
      KS.beta=c(KS.beta$statistic,P.Value=KS.beta$p.value),
      KS.LogNormal=c(KS.LogNormal$statistic,P.Value=KS.LogNormal$p.value),
      KS.Exp=c(KS.Exp$statistic,P.Value=KS.Exp$p.value),
      KS.Normal=c(KS.Normal$statistic,KS.Normal$p.value))
```

```
##           D           P.Value
## KS.Moments 0.05780734 0.3736786147
## KS.beta    0.05239816 0.4985637180
## KS.LogNormal 0.39524722 0.0000000000
## KS.Exp     0.11523305 0.0026158124
## KS.Normal  0.13260203 0.0003039941
```

```
write.csv(One.Minute.Counts.Temps,file="OneMinuteCountsTemps.csv",row.names=FALSE)
Part2.Data<-read.csv("OneMinuteCountsTemps.csv")
head(Part2.Data)
```


```
##      Minute.times Minute.counts Miunute.Temps
## 1           30           7      91.59307
## 2           90          10      97.30860
## 3          150           7      95.98865
## 4          210           4     100.38440
## 5          270           1      99.98330
## 6          330           6     102.54126
```

```
Part2.Data<-Part2.Data[complete.cases(Part2.Data),]
```

```
Part2.Data<-as.data.frame(cbind(Part2.Data,Part2.Data[,2]/60))
colnames(Part2.Data)<-c("Times","Counts","Temperatures","Intensities")
head(Part2.Data)
```

```
##      Times Counts Temperatures Intensities
## 1      30      7      91.59307  0.11666667
## 2      90     10      97.30860  0.16666667
## 3     150      7      95.98865  0.11666667
## 4     210      4     100.38440  0.06666667
## 5     270      1      99.98330  0.01666667
## 6     330      6     102.54126  0.10000000
```


```
plot(Part2.Data$Temperatures,Part2.Data$Intensities)
```



Interpret the plot. What type of

relationship do you observe? There initially seems to be a high amount of data points in the temperatures between 95 to 105 where intensity is between 0 and .3.

```
plot(rank(Part2.Data$Temperatures),rank(Part2.Data$Intensities))
```




What type of dependency you

see in the empirical copula? It seems as temperature increases, so does intensity but there data does seem to be quite sparsed.

```
hist(Part2.Data$Temperatures)
```

Histogram of Part2.Data\$Temperatures



```
head(Part2.Data)
```

```
##      Times Counts Temperatures Intensities
## 1      30       7    91.59307  0.11666667
## 2      90      10    97.30860  0.16666667
## 3     150       7    95.98865  0.11666667
## 4     210       4   100.38440  0.06666667
## 5     270       1    99.98330  0.01666667
## 6     330       6   102.54126  0.10000000
```

```
Part2.Norm <- fitdistr(Part2.Data$Temperatures, "normal")
Part2.Norm$n
```

```
## [1] 242
```

```
ks.test(Part2.Data$Temperatures, "pnorm", mean=100.069853, sd=4.812484)
```

```
##  
## One-sample Kolmogorov-Smirnov test  
##  
## data: Part2.Data$Temperatures  
## D = 0.048912, p-value = 0.6089  
## alternative hypothesis: two-sided
```

data intensities


```
Fitting.Copula <- cbind(pobs(Part2.Data$Temperatures, ties.method = "average"), pobs(Part2.Data$Intensities,  
ties.method = "average"))  
Gumbelcopula<-gumbelCopula(dim=2)  
Copula.Fit<-fitCopula(Gumbelcopula,  
                      pobs(Fitting.Copula, ties.method = "average"),  
                      method = "ml",  
                      optim.method = "BFGS",  
                      optim.control = list(maxit=1000))  
  
summary(Copula.Fit)
```

```
## Call: fitCopula(copula, data = data, method = "ml", optim.method = "BFGS",  
## optim.control = ..3)  
## Fit based on "maximum likelihood" and 242 2-dimensional observations.  
## Gumbel copula, dim. d = 2  
## Estimate Std. Error  
## alpha 1.877 0.099  
## The maximized loglikelihood is 74.18  
## Optimization converged  
## Number of loglikelihood evaluations:  
## function gradient  
## 11 4
```

```
Copula.Fit@estimate
```

```
## [1] 1.877455
```

```
plot(Part2.Data$Intensities)
```



<https://stats.stackexchange.com/questions/90729/generating-values-from-copula-using-copula-package-in-r>

```
Fitting.Copula <- cbind(pobs(Part2.Data$Temperatures, ties.method = "average"), pobs(Part2.Data$Intensities,
ties.method = "average"))
Gumbelcopula<-gumbelCopula(dim=2)
(Copula.Fit<-fitCopula(Gumbelcopula,
                        pobs(Fitting.Copula,ties.method = "average"),
                        method = "ml",
                        optim.method = "BFGS",
                        optim.control = list(maxit=1000)))
```

```
## Call: fitCopula(copula, data = data, method = "ml", optim.method = "BFGS",
##   optim.control = .3)
## Fit based on "maximum likelihood" and 242 2-dimensional observations.
## Copula: gumbelCopula
## alpha
## 1.877
## The maximized loglikelihood is 74.18
## Optimization converged
```

```
summary(Copula.Fit)
```

```
## Call: fitCopula(copula, data = data, method = "ml", optim.method = "BFGS",
##   optim.control = .3)
## Fit based on "maximum likelihood" and 242 2-dimensional observations.
## Gumbel copula, dim. d = 2
##   Estimate Std. Error
## alpha    1.877      0.099
## The maximized loglikelihood is 74.18
## Optimization converged
## Number of loglikelihood evaluations:
## function gradient
##      11          4
```

```
set.seed(8301735)
Simulated.Copula <- as.data.frame(rCopula(copula=gumbelCopula(Copula.Fit@estimate, dim=2), n=250))
Simulated.Copula
```

```
##           V1           V2
## 1  0.346294799 0.279039922
## 2  0.997172131 0.992229381
## 3  0.617043299 0.019520077
```


## 4	0.259385911	0.650830834
## 5	0.270713152	0.027871947
## 6	0.889982559	0.571358980
## 7	0.066110585	0.182890724
## 8	0.107686727	0.519409285
## 9	0.174084135	0.057772586
## 10	0.883338242	0.986061524
## 11	0.484448369	0.161447133
## 12	0.036373304	0.441940142
## 13	0.071734919	0.041168541
## 14	0.553966883	0.734009767
## 15	0.911372278	0.896681713
## 16	0.451324353	0.456238931
## 17	0.260699761	0.488324668
## 18	0.069794134	0.235534249
## 19	0.034087851	0.056572140
## 20	0.038971572	0.179960052
## 21	0.598372551	0.830605495
## 22	0.317778308	0.893942225
## 23	0.888525850	0.772224198
## 24	0.876170025	0.757074080
## 25	0.107540518	0.168576621
## 26	0.733493823	0.772071733
## 27	0.997287721	0.989353221
## 28	0.301781325	0.082714274
## 29	0.954254848	0.962999037
## 30	0.285038849	0.165979563
## 31	0.403090345	0.076256378
## 32	0.645124729	0.160697826
## 33	0.105289411	0.160321658
## 34	0.132595994	0.494041836
## 35	0.317081838	0.131602818
## 36	0.137762855	0.287327393
## 37	0.443768165	0.056579028
## 38	0.949075426	0.611099492
## 39	0.268700433	0.667986777
## 40	0.760405095	0.789516579
## 41	0.251139813	0.344854323
## 42	0.367915345	0.308020507
## 43	0.454789868	0.874486316
## 44	0.342607238	0.240798234
## 45	0.568934380	0.104565289
## 46	0.640065062	0.259939366
## 47	0.860498970	0.808327958
## 48	0.194798105	0.109639468
## 49	0.777644118	0.913489633
## 50	0.190313540	0.069054858
## 51	0.272439974	0.135127469
## 52	0.652139639	0.361091793
## 53	0.208002619	0.507944713
## 54	0.057171737	0.045306302
## 55	0.380310256	0.194585189
## 56	0.673724252	0.839219835
## 57	0.936417204	0.879231489
## 58	0.453455603	0.321425232
## 59	0.317959580	0.176589624
## 60	0.848136578	0.931204745
## 61	0.513045460	0.540847322
## 62	0.118653793	0.374282271
## 63	0.953245854	0.991251027
## 64	0.998450592	0.997083562
## 65	0.750936291	0.766497903
## 66	0.374496993	0.782559428
## 67	0.325793266	0.410983833
## 68	0.055211675	0.754305158
## 69	0.001389858	0.213389195
## 70	0.032047772	0.014793247
## 71	0.410798854	0.305235863
## 72	0.096852102	0.221267846
## 73	0.869480513	0.811375669
## 74	0.553959193	0.402743532
## 75	0.415478448	0.908425787
## 76	0.658439684	0.795733862

```
## 76 0.000100001 0.100700002
## 77 0.500836043 0.177773716
## 78 0.130732654 0.106418497
## 79 0.728538258 0.852964246
## 80 0.506675652 0.346848585
## 81 0.403298707 0.629401476
## 82 0.832650449 0.076076257
## 83 0.871363018 0.903385116
## 84 0.065991802 0.060336663
## 85 0.445598316 0.879416196
## 86 0.098689577 0.189565939
## 87 0.668126578 0.754923937
## 88 0.458781060 0.145391182
## 89 0.616431715 0.443456491
## 90 0.259097933 0.327504692
## 91 0.730065572 0.163266355
## 92 0.557631152 0.317416078
## 93 0.068532221 0.218659056
## 94 0.115281361 0.172396854
## 95 0.027752579 0.741416460
## 96 0.575933466 0.824566096
## 97 0.156754877 0.106585171
## 98 0.541532786 0.481062727
## 99 0.806167096 0.597349549
## 100 0.402087114 0.553718457
## 101 0.190958069 0.608116278
## 102 0.775272388 0.311561955
## 103 0.604865162 0.743569435
## 104 0.402323055 0.719873189
## 105 0.630530759 0.622330896
## 106 0.286804480 0.124848219
## 107 0.895228878 0.977903124
## 108 0.457403485 0.204075542
## 109 0.019012019 0.113302441
## 110 0.347121181 0.326088664
## 111 0.924273251 0.956132090
## 112 0.302643831 0.015504721
## 113 0.295327831 0.203679201
## 114 0.011634054 0.657250128
## 115 0.038694121 0.144764824
## 116 0.623539379 0.443042042
## 117 0.414168145 0.514396115
## 118 0.258503363 0.828638456
## 119 0.300207878 0.227086371
## 120 0.114565619 0.011527819
## 121 0.688129488 0.282178466
## 122 0.033312174 0.187025380
## 123 0.481654377 0.813641936
## 124 0.001295225 0.007634664
## 125 0.388263415 0.074220854
## 126 0.384896269 0.477730364
## 127 0.512750342 0.095197523
## 128 0.104899905 0.112143560
## 129 0.999663790 0.997337416
## 130 0.424814466 0.208653109
## 131 0.850785189 0.906961675
## 132 0.243503710 0.082212315
## 133 0.332952391 0.501024742
## 134 0.614142163 0.064234089
## 135 0.512008780 0.658044816
## 136 0.726432962 0.727007923
## 137 0.831474741 0.762327439
## 138 0.361176648 0.709568185
## 139 0.818923090 0.513360373
## 140 0.514335895 0.341335508
## 141 0.012815634 0.206161948
## 142 0.714056479 0.752363658
## 143 0.057016525 0.397479355
## 144 0.937796579 0.989328463
## 145 0.656339132 0.446326698
## 146 0.429358808 0.532604335
## 147 0.780332827 0.335831855
## 148 0.909961591 0.428400186
```


149 0.853885117 0.854618306
150 0.125915516 0.335436619
151 0.674414150 0.728580537
152 0.575326605 0.544839683
153 0.492399756 0.761673937
154 0.129200037 0.072709677
155 0.463383877 0.980597997
156 0.558399877 0.574285225
157 0.945130219 0.840834077
158 0.881137270 0.761086282
159 0.486699709 0.821690720
160 0.447624655 0.429804159
161 0.020861733 0.003650346
162 0.364095775 0.519227457
163 0.505845012 0.339425551
164 0.769899018 0.808811371
165 0.904011687 0.924558871
166 0.138711509 0.341455797
167 0.156830423 0.268411997
168 0.522202697 0.506267220
169 0.221645707 0.036108822
170 0.625334606 0.917612661
171 0.266936184 0.317186422
172 0.832821183 0.660347274
173 0.429464012 0.143629820
174 0.605314188 0.649099231
175 0.054387878 0.022649230
176 0.755646040 0.483061137
177 0.695940414 0.696213508
178 0.472080113 0.315531163
179 0.245444329 0.203595434
180 0.047595605 0.366907606
181 0.262282368 0.285670622
182 0.956975703 0.919597270
183 0.588160932 0.673302014
184 0.434829350 0.241365353
185 0.667142100 0.869916425
186 0.847527358 0.809397906
187 0.624674068 0.040164227
188 0.027760090 0.190357201
189 0.688352183 0.431352332
190 0.008062198 0.006384882
191 0.341602206 0.472928550
192 0.277975453 0.483515143
193 0.050369720 0.506341692
194 0.637908158 0.329920228
195 0.936923777 0.922430074
196 0.410699085 0.327373201
197 0.480166262 0.952745735
198 0.380469551 0.602995930
199 0.309092956 0.426867719
200 0.228672082 0.152666056
201 0.645131010 0.161050575
202 0.003933024 0.091629976
203 0.518840875 0.520477303
204 0.472784696 0.081789289
205 0.480197294 0.163450804
206 0.273492228 0.415314266
207 0.569764223 0.060069685
208 0.510323457 0.416085480
209 0.195807477 0.375886621
210 0.197995060 0.114624732
211 0.213154338 0.058979585
212 0.781245760 0.955899395
213 0.650924751 0.693425955
214 0.211928103 0.134100289
215 0.437054261 0.379083302
216 0.520008486 0.247978282
217 0.575416121 0.183024244
218 0.936352507 0.961271328
219 0.413688720 0.189479796
220 0.552612441 0.764121181
221 0.129377091 0.649275765


```
## 222 0.483406409 0.124640055
## 223 0.845764079 0.989894190
## 224 0.257035433 0.191943058
## 225 0.544792866 0.222054567
## 226 0.938533201 0.653926171
## 227 0.284342485 0.471243034
## 228 0.192515105 0.512119067
## 229 0.708638879 0.256203151
## 230 0.644266174 0.691958491
## 231 0.022494432 0.082467557
## 232 0.955267032 0.908984022
## 233 0.583265824 0.508326189
## 234 0.038639628 0.243757367
## 235 0.545428782 0.452583523
## 236 0.211175189 0.098317189
## 237 0.260897780 0.628421656
## 238 0.758170311 0.671253193
## 239 0.455266204 0.401976980
## 240 0.292001213 0.220047480
## 241 0.864043637 0.484360690
## 242 0.610785122 0.127772752
## 243 0.580441478 0.676647074
## 244 0.964204272 0.884034463
## 245 0.485969736 0.661506996
## 246 0.853547470 0.880471408
## 247 0.389515538 0.344212043
## 248 0.160170189 0.126158432
## 249 0.086382786 0.563810495
## 250 0.917273320 0.966115851
```


```
persp(gumbelCopula(Copula.Fit@estimate), dCopula, main="pdf", xlab="u", ylab="v", zlab="c(u,v)")
```




```
contour(gumbelCopula(Copula.Fit@estimate), dCopula, main="pdf", xlab="u", ylab="v")
```




```
plot(Simulated.Copula)
```




```
df <- data.frame(Temperatures = qnorm(Simulated.Copula[,1],mean=100.069853 ,sd=4.812484),
                 Intensities = qgamma(Simulated.Copula[,2],shape=1.655712, rate=8.132182/60 ))
Emperical.Copula <- as.data.frame(apply(df,2,rank))
plot(Emperical.Copula)
```



```
set.seed(8301735)
Simulated.Copula1 <- as.data.frame(rCopula(copula=gumbelCopula(Copula.Fit@estimate,dim=2), n=5000))
df1 <- data.frame(Temperatures = qnorm(Simulated.Copula1[,1],mean=100.069853 ,sd=4.812484),
                  Intensities = qgamma(Simulated.Copula1[,2],shape=1.655712, rate=8.132182 ))
plot(df1)
```



```
Emperical.Copula1 <- as.data.frame(apply(df1,2,rank))
plot(Emperical.Copula1)
```




My numbers are slightly diferent than Yuri's below and I'm not exactly sure as to why. Have mercy! haaaha

```
head(Part2.Data)
```

```
##      Times Counts Temperatures Intensities
## 1      30      7      91.59307  0.11666667
## 2      90     10      97.30860  0.16666667
## 3     150      7      95.98865  0.11666667
## 4     210      4     100.38440  0.06666667
## 5     270      1      99.98330  0.01666667
## 6     330      6     102.54126  0.10000000
```

```
plot(Part2.Data$Temperatures,Part2.Data$Intensities)
```



```
NB.Fit.To.Sample <- glm.nb(Counts~Temperatures,Part2.Data)
NB.Fit.To.Sample$coefficients
```

```
## (Intercept) Temperatures
## -7.43137538 0.09843241
```

```
NB.Fit.To.Sample$deviance
```

```
## [1] 257.1937
```

```
NB.Fit.To.Sample$df.residual
```

```
## [1] 240
```

```
NB.Fit.To.Sample$aic
```


```
## [1] 1557.817
```

```
NB.Fit.To.Sample$theta
```

```
## [1] 4.202612
```

```
Simulated.Temperature <- qnorm(Simulated.Copula[,1],mean=100.069853 ,sd=4.812484)
Simulated.Intensities <- qgamma(Simulated.Copula[,2],shape=1.655712, rate=8.132182 )
Simulated.Tails<-as.data.frame(
  cbind(round(Simulated.Intensities[ (Simulated.Temperature>110) & (Simulated.Intensities>.5) ]*60),
    Simulated.Temperature[ (Simulated.Temperature>110) & (Simulated.Intensities>.5) ]))
colnames(Simulated.Tails)<-c("Counts", "Temperatures")
```

```
plot(Simulated.Tails$Temperatures,Simulated.Tails$Counts)
```



With the Simulated tails, there is a lot less overdispersion which means data towards very high temperatures has a high correlation with the Intensity.

```
NB.Simulated.Tails <- glm.nb(Counts~Temperatures,Simulated.Tails)
```

```
## Warning in theta.ml(Y, mu, sum(w), w, limit = control$maxit, trace = control$trace > : iteration
## limit reached

## Warning in theta.ml(Y, mu, sum(w), w, limit = control$maxit, trace = control$trace > : iteration
## limit reached
```

```
NB.Simulated.Tails$theta
```

```
## [1] 737219.5
```

Is there an alternative model that you would try to fit to the simulated tail data? I would also try fitting a poisson model which we will fit below
What do both models tell you about the relationships between the temperature and the counts? They don't fit well because of the high theta, but the models do a moderately good job of fitting the data.

```
Poisson.Fit <- glm(Counts~Temperatures, Simulated.Tails, family="poisson")
Poisson.Fit$deviance
```

```
## [1] 62.72328
```

```
summary(Poisson.Fit)$df
```

```
## [1] 2 73 2
```

```
Poisson.Fit$aic
```

```
## [1] 490.0239
```

```
dispersiontest(Poisson.Fit)
```

```
##
## Overdispersion test
##
## data: Poisson.Fit
## z = -1.43, p-value = 0.9236
## alternative hypothesis: true dispersion is greater than 1
## sample estimates:
## dispersion
## 0.8305643
```